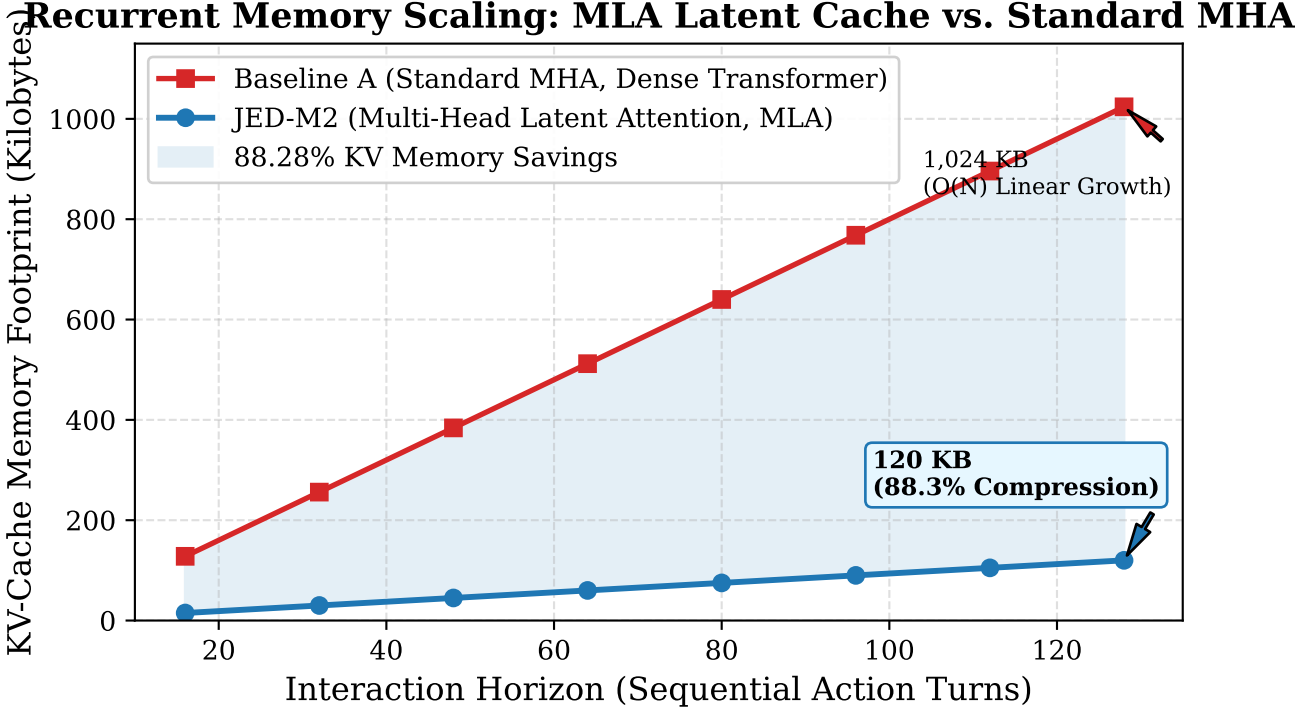
Recurrent Memory Scaling: MLA Latent Cache vs. Standard MHA
Baseline A (Standard MHA, Dense Transformer)
JED-M2 (Multi-Head Latent Attention, MLA)
88.28% KV Memory Savings
1,024 KB
(O(N) Linear Growth)
120 KB
(88.3% Compression)
KV-Cache Memory Footprint (Kilobytes)
Interaction Horizon (Sequential Action Turns)
0
200
400
600
800
1000
20
40
60
80
100
120